

PROJECT OVERVIEW

In this project, I will focus on data cleaning, imputation, analysis, and visualization to derive valuable insights for a business stakeholder.

Business understanding

Our company is expanding in to new industries to diversify its portfolio. Specifically, they are interested in purchasing and operating airplanes for commercial and private enterprises, In order to pursue this venture the company has to be in the know of potential risks of this venture. This includes the number of accidents that each aircraft has had over the years, the market demand and the safety of the aircraft. I have been charged with determining which aircraft are the lowest risk for the company to start this new business endeavor. I will thereby translate my findings into actionable insights that the head of the new aviation division can use to help decide which aircraft to purchase.

The data

In the data folder is a dataset from the National Transportation Safety Board that includes aviation accident data from 1962 to 2023 about civil aviation accidents and selected incidents in the United States and international waters.

1. My analysis would yield three concrete business recommendations
2. Communication
3. Visualization

Import necessary libraries to be used in the project

```
#IMPORTING ALL NECESSARY LIBRARIES.
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
%matplotlib inline

#This part here ignores warnings
from pandas.errors import DtypeWarning

warnings.filterwarnings("ignore", category=DtypeWarning)
```

Load the dataset

```
#This reads the dataset that I want to work on
df = pd.read_csv("AviationData.csv",encoding = 'Windows-1252')
```

Data sanity check

```
#This gives us the shape of the dataset.It has (88889)rows and
(31)columns
```

```
df.shape
```

```
(88889, 31)
```

```
#This give an outline of the first 5 rows of the dataset
```

```
df.head()
```

	Event.Id	Investigation.Type	Accident.Number	Event.Date	\
0	20001218X45444	Accident	SEA87LA080	1948-10-24	
1	20001218X45447	Accident	LAX94LA336	1962-07-19	
2	20061025X01555	Accident	NYC07LA005	1974-08-30	
3	20001218X45448	Accident	LAX96LA321	1977-06-19	
4	20041105X01764	Accident	CHI79FA064	1979-08-02	

	Location	Country	Latitude	Longitude	Airport.Code	\
0	MOOSE CREEK, ID	United States	NaN	NaN	NaN	
1	BRIDGEPORT, CA	United States	NaN	NaN	NaN	
2	Saltville, VA	United States	36.9222	-81.8781	NaN	
3	EUREKA, CA	United States	NaN	NaN	NaN	
4	Canton, OH	United States	NaN	NaN	NaN	

	Airport.Name	...	Purpose.of.flight	Air.carrier	Total.Fatal.Injuries	\
0	NaN	...	Personal	NaN	2.0	
1	NaN	...	Personal	NaN	4.0	
2	NaN	...	Personal	NaN	3.0	
3	NaN	...	Personal	NaN	2.0	
4	NaN	...	Personal	NaN	1.0	

	Total.Serious.Injuries	Total.Minor.Injuries	Total.Uninjured	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	NaN	NaN	NaN	
3	0.0	0.0	0.0	
4	2.0	NaN	0.0	

Weather.Condition	Broad.phase.of.flight	Report.Status
Publication.Date		
0	UNK	Cruise Probable Cause
NaN		
1	UNK	Unknown Probable Cause 19-
09-1996		
2	IMC	Cruise Probable Cause 26-
02-2007		
3	IMC	Cruise Probable Cause 12-
09-2000		
4	VMC	Approach Probable Cause 16-
04-1980		

[5 rows x 31 columns]

#This give an outline of the last 5 rows of the dataset
df.tail()

Event.Id	Investigation.Type	Accident.Number
Event.Date \		
88884 20221227106491	Accident	ERA23LA093 2022-12-26
88885 20221227106494	Accident	ERA23LA095 2022-12-26
88886 20221227106497	Accident	WPR23LA075 2022-12-26
88887 20221227106498	Accident	WPR23LA076 2022-12-26
88888 20221230106513	Accident	ERA23LA097 2022-12-29

Location	Country	Latitude	Longitude	Airport.Code
88884 Annapolis, MD	United States	NaN	NaN	NaN
88885 Hampton, NH	United States	NaN	NaN	NaN
88886 Payson, AZ	United States	341525N	1112021W	PAN
88887 Morgan, UT	United States	NaN	NaN	NaN
88888 Athens, GA	United States	NaN	NaN	NaN

Airport.Name	...	Purpose.of.flight	Air.carrier
88884 NaN	...	Personal	NaN
88885 NaN	...	NaN	NaN
88886 PAYSON	...	Personal	NaN
88887 NaN	...	Personal	MC CESSNA 210N LLC
88888 NaN	...	Personal	NaN

Total.Fatal.Injuries	Total.Serious.Injuries	Total.Minor.Injuries
\		
88884	0.0	1.0 0.0
88885	0.0	0.0 0.0

88886	0.0	0.0	0.0
88887	0.0	0.0	0.0
88888	0.0	1.0	0.0

	Total.Uninjured	Weather.Condition	Broad.phase.of.flight
Report.Status \			
88884	0.0	NaN	NaN
NaN			
88885	0.0	NaN	NaN
NaN			
88886	1.0	VMC	NaN
NaN			
88887	0.0	NaN	NaN
NaN			
88888	1.0	NaN	NaN
NaN			

	Publication.Date
88884	29-12-2022
88885	NaN
88886	27-12-2022
88887	NaN
88888	30-12-2022

[5 rows x 31 columns]

#Gives a look into the columns inside the dataframe

df.columns

```
Index(['Event.Id', 'Investigation.Type', 'Accident.Number',
      'Event.Date',
      'Location', 'Country', 'Latitude', 'Longitude', 'Airport.Code',
      'Airport.Name', 'Injury.Severity', 'Aircraft.damage',
      'Aircraft.Category', 'Registration.Number', 'Make', 'Model',
      'Amateur.Built', 'Number.ofEngines', 'Engine.Type',
      'FAR.Description',
      'Schedule', 'Purpose.of.flight', 'Air.carrier',
      'Total.Fatal.Injuries',
      'Total.Serious.Injuries', 'Total.Minor.Injuries',
      'Total.Uninjured',
      'Weather.Condition', 'Broad.phase.of.flight', 'Report.Status',
      'Publication.Date'],
      dtype='object')
```

#This gives us the summary statistics of the dataset

df.describe()

	Number.of.Engines	Total.Fatal.Injuries	Total.Serious.Injuries
\			
count	82805.000000	77488.000000	76379.000000
mean	1.146585	0.647855	0.279881
std	0.446510	5.485960	1.544084
min	0.000000	0.000000	0.000000
25%	1.000000	0.000000	0.000000
50%	1.000000	0.000000	0.000000
75%	1.000000	0.000000	0.000000
max	8.000000	349.000000	161.000000

	Total.Minor.Injuries	Total.Uninjured
count	76956.000000	82977.000000
mean	0.357061	5.325440
std	2.235625	27.913634
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	1.000000
75%	0.000000	2.000000
max	380.000000	699.000000

#This give summary statistics of objects

```
df.describe(include="object")
```


	Event.Id	Investigation.Type	Accident.Number	Event.Date
\				
count	88889	88889	88889	88889
unique	87951	2	88863	14782
top	20001214X45071	Accident	ERA22LA379	1982-05-16
freq	3	85015	2	25

	Location	Country	Latitude	Longitude
Airport.Code				
\				
count	88837	88663	34382	34373
unique	27758	219	25592	27156
top	ANCHORAGE, AK	United States	332739N	0112457W

freq	434	82248	19	24	1488
	Airport.Name	... Amateur.Built	Engine.Type	FAR.Description	
\	count	52790	88787	81812	32023
unique	24871	...	2	13	31
top	Private	...	No	Reciprocating	091
freq	240	...	80312	69530	18221

	Schedule	Purpose.of.flight	Air.carrier	Weather.Condition	\
count	12582	82697	16648	84397	
unique	3	26	13590	4	
top	NSCH	Personal	Pilot	VMC	
freq	4474	49448	258	77303	

	Broad.phase.of.flight	Report.Status	Publication.Date
count	61724	82508	75118
unique	12	17075	2924
top	Landing	Probable Cause	25-09-2020
freq	15428	61754	17019

[4 rows x 26 columns]

#This gives the total overview of our dataset
#From the look of the data, we have a couple of missing data
df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 88889 entries, 0 to 88888
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Event.Id                             88889 non-null  object
1   Investigation.Type                    88889 non-null  object
2   Accident.Number                      88889 non-null  object
3   Event.Date                           88889 non-null  object
4   Location                             88837 non-null  object
5   Country                             88663 non-null  object
6   Latitude                             34382 non-null  object
7   Longitude                            34373 non-null  object
8   Airport.Code                         50249 non-null  object
9   Airport.Name                         52790 non-null  object
10  Injury.Severity                      87889 non-null  object
11  Aircraft.damage                      85695 non-null  object
12  Aircraft.Category                    32287 non-null  object
13  Registration.Number                  87572 non-null  object
```

14	Make	88826	non-null	object
15	Model	88797	non-null	object
16	Amateur.Built	88787	non-null	object
17	Number.of.Engines	82805	non-null	float64
18	Engine.Type	81812	non-null	object
19	FAR.Description	32023	non-null	object
20	Schedule	12582	non-null	object
21	Purpose.of.flight	82697	non-null	object
22	Air.carrier	16648	non-null	object
23	Total.Fatal.Injuries	77488	non-null	float64
24	Total.Serious.Injuries	76379	non-null	float64
25	Total.Minor.Injuries	76956	non-null	float64
26	Total.Uninjured	82977	non-null	float64
27	Weather.Condition	84397	non-null	object
28	Broad.phase.of.flight	61724	non-null	object
29	Report.Status	82508	non-null	object
30	Publication.Date	75118	non-null	object

dtypes: float64(5), object(26)

memory usage: 21.0+ MB

#Now lets find the missing values

#Well,we have a number of collumns that have missing values

df.isna().sum()

Event.Id	0
Investigation.Type	0
Accident.Number	0
Event.Date	0
Location	52
Country	226
Latitude	54507
Longitude	54516
Airport.Code	38640
Airport.Name	36099
Injury.Severity	1000
Aircraft.damage	3194
Aircraft.Category	56602
Registration.Number	1317
Make	63
Model	92
Amateur.Built	102
Number.of.Engines	6084
Engine.Type	7077
FAR.Description	56866
Schedule	76307
Purpose.of.flight	6192
Air.carrier	72241
Total.Fatal.Injuries	11401
Total.Serious.Injuries	12510
Total.Minor.Injuries	11933

Total.Uninjured	5912
Weather.Condition	4492
Broad.phase.of.flight	27165
Report.Status	6381
Publication.Date	13771

dtype: int64

#Lets now check the percentage of missing values in the columns so as it can give us insights on what columns we would drop

```
df.isna().sum()/len(df)*100
```

Event.Id	0.000000
Investigation.Type	0.000000
Accident.Number	0.000000
Event.Date	0.000000
Location	0.058500
Country	0.254250
Latitude	61.320298
Longitude	61.330423
Airport.Code	43.469946
Airport.Name	40.611324
Injury.Severity	1.124999
Aircraft.damage	3.593246
Aircraft.Category	63.677170
Registration.Number	1.481623
Make	0.070875
Model	0.103500
Amateur.Built	0.114750
Number.ofEngines	6.844491
Engine.Type	7.961615
FAR.Description	63.974170
Schedule	85.845268
Purpose.of.flight	6.965991
Air.carrier	81.271023
Total.Fatal.Injuries	12.826109
Total.Serious.Injuries	14.073732
Total.Minor.Injuries	13.424608
Total.Uninjured	6.650992
Weather.Condition	5.053494
Broad.phase.of.flight	30.560587
Report.Status	7.178616
Publication.Date	15.492356

dtype: float64

Great, We now know what columns have massive numbers of missing data and dropping them will be necessary. Not dropping them would affect what I really want to achieve.

Cleaning the dataset

dropping the columns that won't be needed. All the columns that I would be dropping I will be storing them in the columns_dropped variable.

```
#placing unnecessary columns in a variable
columns_dropped = ['Accident.Number', 'Location', 'Country', 'Latitude',
'Longitude', 'Airport.Code', 'Airport.Name', 'Registration.Number',
'FAR.Description',
'Schedule', 'Air.carrier', 'Publication.Date']
```

```
#dropping the columns from the original dataframe to refined one
cleaned_df = df.drop(columns=columns_dropped)
```

```
#first five rows of the dataset
cleaned_df.head()
```

	Event.Id	Investigation.Type	Event.Date	Injury.Severity	\
0	20001218X45444	Accident	1948-10-24	Fatal(2)	
1	20001218X45447	Accident	1962-07-19	Fatal(4)	
2	20061025X01555	Accident	1974-08-30	Fatal(3)	
3	20001218X45448	Accident	1977-06-19	Fatal(2)	
4	20041105X01764	Accident	1979-08-02	Fatal(1)	

	Aircraft.damage	Aircraft.Category	Make	Model	Amateur.Built	\
0	Destroyed	NaN	Stinson	108-3	No	
1	Destroyed	NaN	Piper	PA24-180	No	
2	Destroyed	NaN	Cessna	172M	No	
3	Destroyed	NaN	Rockwell	112	No	
4	Destroyed	NaN	Cessna	501	No	

	Number.of.Engines	Engine.Type	Purpose.of.flight	Total.Fatal.Injuries	\
0	1.0	Reciprocating	Personal	2.0	
1	1.0	Reciprocating	Personal	4.0	
2	1.0	Reciprocating	Personal	3.0	
3	1.0	Reciprocating	Personal	2.0	
4	NaN	NaN	Personal	1.0	

	Total.Serious.Injuries	Total.Minor.Injuries	Total.Uninjured	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	NaN	NaN	NaN	
3	0.0	0.0	0.0	
4	2.0	NaN	0.0	

	Weather.Condition	Broad.phase.of.flight	Report.Status
0	UNK	Cruise	Probable Cause
1	UNK	Unknown	Probable Cause
2	IMC	Cruise	Probable Cause
3	IMC	Cruise	Probable Cause
4	VMC	Approach	Probable Cause

#finding duplicates

```
cleaned_df.duplicated()
```

```
0      False
1      False
2      False
3      False
4      False
```

```
...
88884   False
88885   False
88886   False
88887   False
88888   False
```

```
Length: 88889, dtype: bool
```

#Checking out the duplicates and printing them out

```
duplicates = cleaned_df[df.duplicated()]
```

```
print(duplicates)
```

```
Empty DataFrame
```

```
Columns: [Event.Id, Investigation.Type, Event.Date, Injury.Severity,
Aircraft.damage, Aircraft.Category, Make, Model, Amateur.Built,
Number.ofEngines, Engine.Type, Purpose.of.flight,
Total.Fatal.Injuries, Total.Serious.Injuries, Total.Minor.Injuries,
Total.Uninjured, Weather.Condition, Broad.phase.of.flight,
Report.Status]
Index: []
```

#Choosing a primary key which helps in removing duplicates

```
cleaned_df = cleaned_df.drop_duplicates(subset="Event.Id")
```

#Now lets do a quick check at our dataframe to see if there are duplicates

```
duplicate_count = cleaned_df.duplicated().sum()
```

```
duplicate_count
print(f"Number of duplicate rows: {duplicate_count}")
```

Number of duplicate rows: 0

Great, We dropped the duplicates.

Data Analysis

```
#This line basically ignores the warnings
import warnings
warnings.filterwarnings('ignore')
#This defines a function called aircraft_category and identifies the
Nan values and replaces them with Unknown
def aircraft_category(value):
    if pd.isna(value) or value.strip() == "":
        return "Unknown"
    return value.strip().title()

# Apply it to the column
cleaned_df["Aircraft.Category"] =
cleaned_df["Aircraft.Category"].apply(aircraft_category)

#to check on the first 5 rows
cleaned_df.head()
```

	Event.Id	Investigation.Type	Event.Date	Injury.Severity	\
0	20001218X45444	Accident	1948-10-24	Fatal(2)	
1	20001218X45447	Accident	1962-07-19	Fatal(4)	
2	20061025X01555	Accident	1974-08-30	Fatal(3)	
3	20001218X45448	Accident	1977-06-19	Fatal(2)	
4	20041105X01764	Accident	1979-08-02	Fatal(1)	

	Aircraft.damage	Aircraft.Category	Make	Model	Amateur.Built	\
0	Destroyed	Unknown	Stinson	108-3	No	
1	Destroyed	Unknown	Piper	PA24-180	No	
2	Destroyed	Unknown	Cessna	172M	No	
3	Destroyed	Unknown	Rockwell	112	No	
4	Destroyed	Unknown	Cessna	501	No	

	Number.of.Engines	Engine.Type	Purpose.of.flight
Total.Fatal.Injuries	\		
0	1.0	Reciprocating	Personal

2.0				
1	1.0	Reciprocating	Personal	
4.0				
2	1.0	Reciprocating	Personal	
3.0				
3	1.0	Reciprocating	Personal	
2.0				
4	NaN	NaN	Personal	
1.0				
	Total.Serious.Injuries	Total.Minor.Injuries	Total.Uninjured	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	NaN	NaN	NaN	
3	0.0	0.0	0.0	
4	2.0	NaN	0.0	
	Weather.Condition	Broad.phase.of.flight	Report.Status	
0	UNK	Cruise	Probable Cause	
1	UNK	Unknown	Probable Cause	
2	IMC	Cruise	Probable Cause	
3	IMC	Cruise	Probable Cause	
4	VMC	Approach	Probable Cause	

Now, lets check for the other existing nan values

```
#Checks for the number of nan values in each and every column in my
cleaned_df
cleaned_df.isna().sum()

Event.Id                0
Investigation.Type      0
Event.Date              0
Injury.Severity         990
Aircraft.damage        3103
Aircraft.Category       0
Make                   63
Model                  92
Amateur.Built          100
Number.ofEngines       6027
Engine.Type            7024
Purpose.of.flight      6122
Total.Fatal.Injuries   11267
Total.Serious.Injuries 12322
Total.Minor.Injuries   11760
Total.Uninjured        5863
Weather.Condition      4473
Broad.phase.of.flight  27114
```

```
Report.Status          6361
dtype: int64
```

```
#This if and ifelse loop checks out our data set and replaces the nan values with median for numerical values and mode for categorical values
```

```
for column in cleaned_df.columns:
    if cleaned_df[column].dtype == 'object':
        #Categorical value filled with the most frequent value(mode)
        cleaned_df[column] =
cleaned_df[column].fillna(cleaned_df[column].mode()[0])
    else:
        #Numerical value filled with median
        cleaned_df[column] =
cleaned_df[column].fillna(cleaned_df[column].median())
cleaned_df.head()
```

	Event.Id	Investigation.Type	Event.Date	Injury.Severity	\
0	20001218X45444	Accident	1948-10-24	Fatal(2)	
1	20001218X45447	Accident	1962-07-19	Fatal(4)	
2	20061025X01555	Accident	1974-08-30	Fatal(3)	
3	20001218X45448	Accident	1977-06-19	Fatal(2)	
4	20041105X01764	Accident	1979-08-02	Fatal(1)	

	Aircraft.damage	Aircraft.Category	Make	Model	Amateur.Built
0	Destroyed	Unknown	Stinson	108-3	No
1	Destroyed	Unknown	Piper	PA24-180	No
2	Destroyed	Unknown	Cessna	172M	No
3	Destroyed	Unknown	Rockwell	112	No
4	Destroyed	Unknown	Cessna	501	No

	Number.of.Engines	Engine.Type	Purpose.of.flight
Total.Fatal.Injuries			
0	1.0	Reciprocating	Personal
2.0			
1	1.0	Reciprocating	Personal
4.0			
2	1.0	Reciprocating	Personal
3.0			
3	1.0	Reciprocating	Personal
2.0			
4	1.0	Reciprocating	Personal
1.0			

	Total.Serious.Injuries	Total.Minor.Injuries	Total.Uninjured	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	1.0	
3	0.0	0.0	0.0	
4	2.0	0.0	0.0	

	Weather.Condition	Broad.phase.of.flight	Report.Status
0	UNK	Cruise	Probable Cause
1	UNK	Unknown	Probable Cause
2	IMC	Cruise	Probable Cause
3	IMC	Cruise	Probable Cause
4	VMC	Approach	Probable Cause

#Rechecking if there are any nan values

```
cleaned_df.isna().sum()
```

```
Event.Id          0
Investigation.Type 0
Event.Date        0
Injury.Severity   0
Aircraft.damage   0
Aircraft.Category 0
Make              0
Model             0
Amateur.Built     0
Number.of.Engines 0
Engine.Type       0
Purpose.of.flight 0
Total.Fatal.Injuries 0
Total.Serious.Injuries 0
Total.Minor.Injuries 0
Total.Uninjured   0
Weather.Condition 0
Broad.phase.of.flight 0
Report.Status     0
dtype: int64
```

Great! the dataset does not nan values

```
import warnings
warnings.filterwarnings('ignore')
#The is a little bit of differences in values in my dataset. look at the Make column.Cessna tends to be the same as CESSNA.
#To fix this, I am making all the values in this column in capital letters
cleaned_df['Make'] = cleaned_df['Make'].str.upper()

cleaned_df.head()
```

	Event.Id	Investigation.Type	Event.Date	Injury.Severity	\
0	20001218X45444	Accident	1948-10-24	Fatal(2)	
1	20001218X45447	Accident	1962-07-19	Fatal(4)	
2	20061025X01555	Accident	1974-08-30	Fatal(3)	
3	20001218X45448	Accident	1977-06-19	Fatal(2)	
4	20041105X01764	Accident	1979-08-02	Fatal(1)	

	Aircraft.damage	Aircraft.Category	Make	Model	Amateur.Built	\
0	Destroyed	Unknown	STINSON	108-3	No	
1	Destroyed	Unknown	PIPER	PA24-180	No	
2	Destroyed	Unknown	CESSNA	172M	No	
3	Destroyed	Unknown	ROCKWELL	112	No	
4	Destroyed	Unknown	CESSNA	501	No	

	Number.of.Engines	Engine.Type	Purpose.of.flight
Total.Fatal.Injuries			
0	1.0	Reciprocating	Personal
2.0			
1	1.0	Reciprocating	Personal
4.0			
2	1.0	Reciprocating	Personal
3.0			
3	1.0	Reciprocating	Personal
2.0			
4	1.0	Reciprocating	Personal
1.0			

	Total.Serious.Injuries	Total.Minor.Injuries	Total.Uninjured	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	1.0	
3	0.0	0.0	0.0	
4	2.0	0.0	0.0	

	Weather.Condition	Broad.phase.of.flight	Report.Status
0	UNK	Cruise	Probable Cause
1	UNK	Unknown	Probable Cause
2	IMC	Cruise	Probable Cause
3	IMC	Cruise	Probable Cause
4	VMC	Approach	Probable Cause

cleaned_df.shape

(87951, 19)

Great, the dataset is now clean.

Visualizations

Great, now that we have a much clearer dataset, let us now visualise it to make a more informed decision. First, let us plot a bargraph to see the number of accidents of all the aircrafts in our dataset over the years to see which aircraft has the most number of accidents.

```
#creates a new dataframe called Event.Date and uses pandas to read it as a date
cleaned_df['Event.Date'] = pd.to_datetime(cleaned_df['Event.Date'],
errors='coerce')

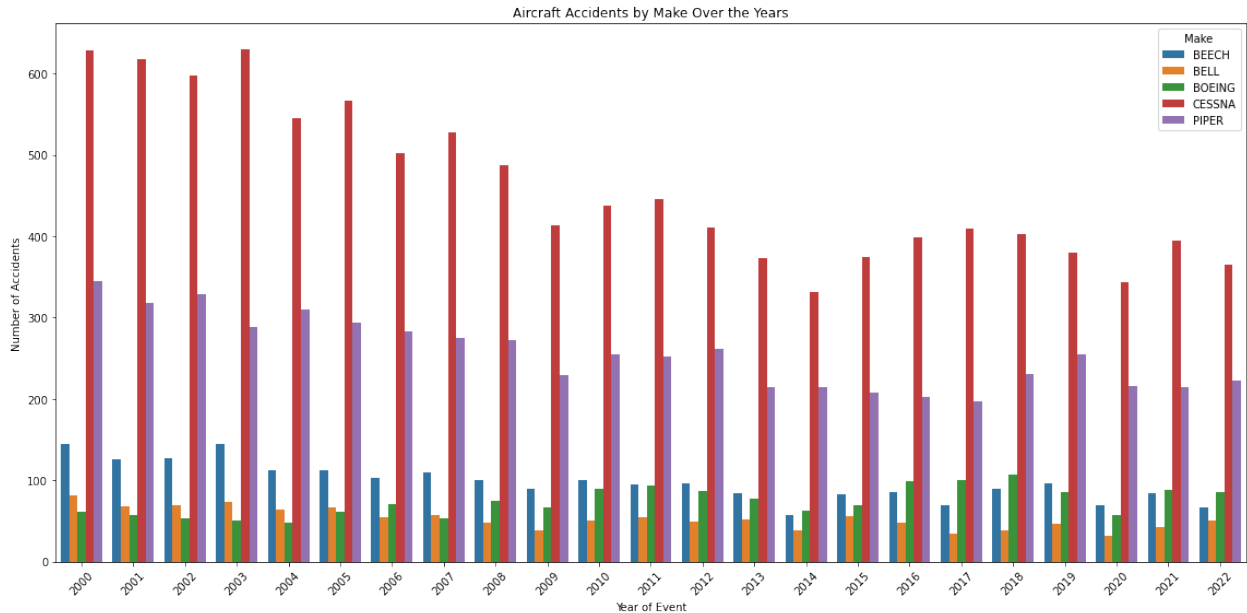
# Create Event.Year column
cleaned_df['Event.Year'] = cleaned_df['Event.Date'].dt.year
#Creates a new column that only contains the year part.This is why we are using dt.year in my previous code.
grouped = cleaned_df.groupby(['Event.Year',
'Make']).size().reset_index(name='Accident Count')

# Optional: Filter recent years and top makes
grouped = grouped[grouped['Event.Year'] >= 2000]
top_makes = cleaned_df['Make'].value_counts().head(5).index
grouped = grouped[grouped['Make'].isin(top_makes)]

# Plot
plt.figure(figsize=(16,8))
ax = sns.barplot(data=grouped, x='Event.Year', y='Accident Count',
hue='Make', dodge=True, palette='tab10')

#Customize
ax.set_title('Aircraft Accidents by Make Over the Years')
ax.set_xlabel('Year of Event')
ax.set_ylabel('Number of Accidents')
ax.legend(title='Make', loc='upper right')
plt.xticks(rotation=45)
plt.tight_layout()

#saves image to my device
plt.savefig('Aircraft Accidents by Make Over the Years.png', dpi=300)
plt.show()
```

oh, Perfect we now know basing on this bargraph that Cessna has been having more accidents over the years.

Now,lets have a look at each make and the number of events based on the event ID.

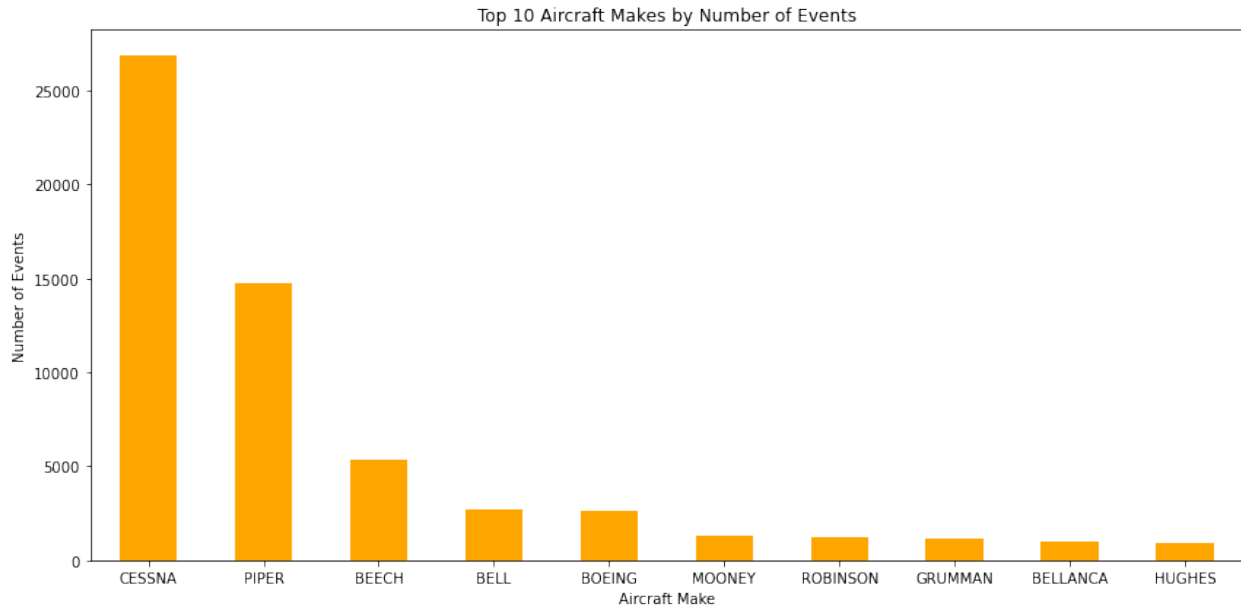
```
# Group by Make and count
event_counts = cleaned_df.groupby('Make')
['Event.Id'].nunique().sort_values(ascending=False)

# Only take top 10
top_10_makes = event_counts.head(10)

# Plot
plt.figure(figsize=(12,6))
top_10_makes.plot(kind='bar', color='orange')

plt.title('Top 10 Aircraft Makes by Number of Events')
plt.xlabel('Aircraft Make')
plt.ylabel('Number of Events')
plt.xticks(rotation=360)

#saves image to my device
plt.savefig('Top 10 Aircraft Makes by Number of Events', dpi=300)
plt.tight_layout()
```



Again, CESSNA seems to have a huge number of events compared to the rest of the aircrafts.

Lets take a look at CESSNA

```
# Filter data for CESSNA models
cessna_data = cleaned_df[cleaned_df['Make'] == 'CESSNA']

# Calculate the number of accidents for each Cessna model
accidents_by_model = cessna_data['Model'].value_counts().head(10)

# Calculate the total number of fatal injuries for each Cessna model
fatalities_by_model = cessna_data.groupby('Model')
['Total.Fatal.Injuries'].sum()


# Calculate the fatality rate (total fatal injuries / number of
# accidents). multiplying it by 100 turns it to a percentage.
fatality_rate = (fatalities_by_model /
accidents_by_model).dropna().sort_values(ascending=False) * 100

# Plot the fatality rates for Cessna models
plt.figure(figsize=(14, 7))
plt.title('Fatality Rates by Cessna Aircraft Model', fontsize=14)
fatality_rate.plot(kind='bar', color='orange', width=0.4)
plt.xlabel('Cessna Model', fontsize=12)
plt.ylabel('Fatality Rate (%)', fontsize=12)
plt.xticks(rotation=360)

# Adjust layout to make room for labels
plt.tight_layout()

#saves image to my device
```

```
plt.savefig('Fatality Rates by Cessna Aircraft Model', dpi=300)
plt.show()
```



The most predominant model is the cessna 172N thereby making it the most used. Lets get a dive into BOEING.

```
# Filter data for Boeing models
boeing_data = cleaned_df[cleaned_df['Make'] == 'BOEING']

# Calculate the number of accidents for each Cessna model
accidents_by_model = boeing_data['Model'].value_counts().head(10)

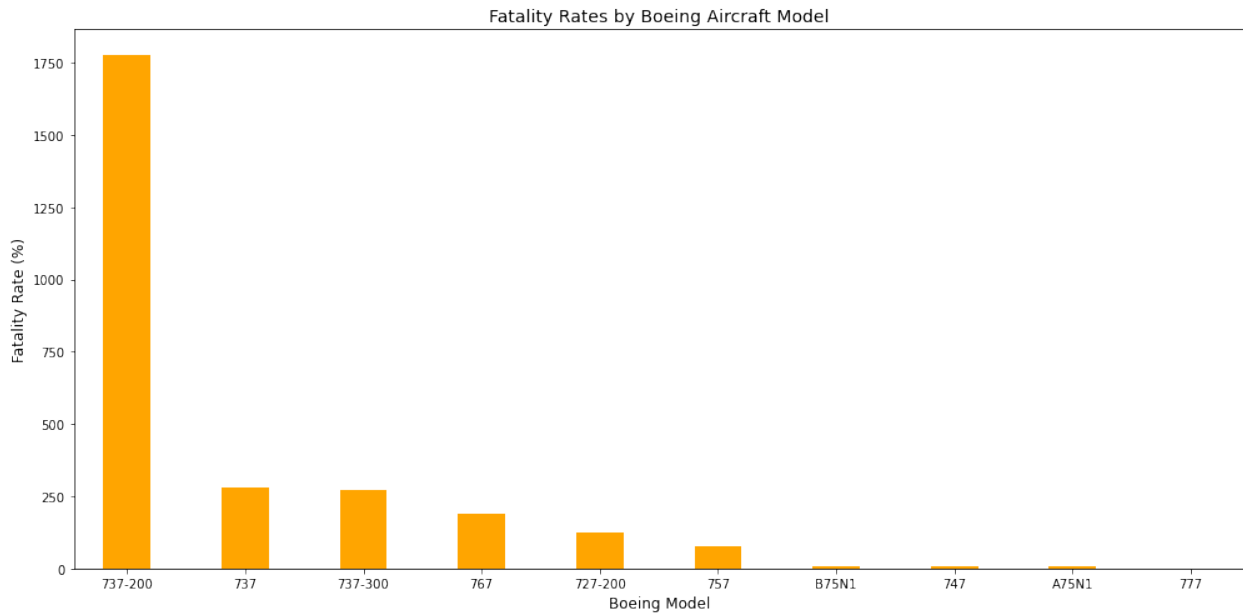
# Calculate the total number of fatal injuries for each Cessna model
fatalities_by_model = boeing_data.groupby('Model')
['Total.Fatal.Injuries'].sum()

# Calculate the fatality rate (total fatal injuries / number of
# accidents). multiplying it by 100 turns it to a percentage.
fatality_rate = (fatalities_by_model /
accidents_by_model).dropna().sort_values(ascending=False) * 100

# Plot the fatality rates for Cessna models
plt.figure(figsize=(14, 7))
plt.title('Fatality Rates by Boeing Aircraft Model', fontsize=14)
fatality_rate.plot(kind='bar', color='orange', width=0.4)
plt.xlabel('Boeing Model', fontsize=12)
plt.ylabel('Fatality Rate (%)', fontsize=12)
plt.xticks(rotation=360)

# Adjust layout to make room for labels
plt.tight_layout()
```

```
#saves image to my device
plt.savefig('Fatality Rates by boeing Aircraft Mode', dpi=300)
plt.show()
```



BOEING 737-200 seems to be the most predominant model based on the fatality rate thereby making it the most used model.

Now, lets check the number of accidents basing on the phase of flight.

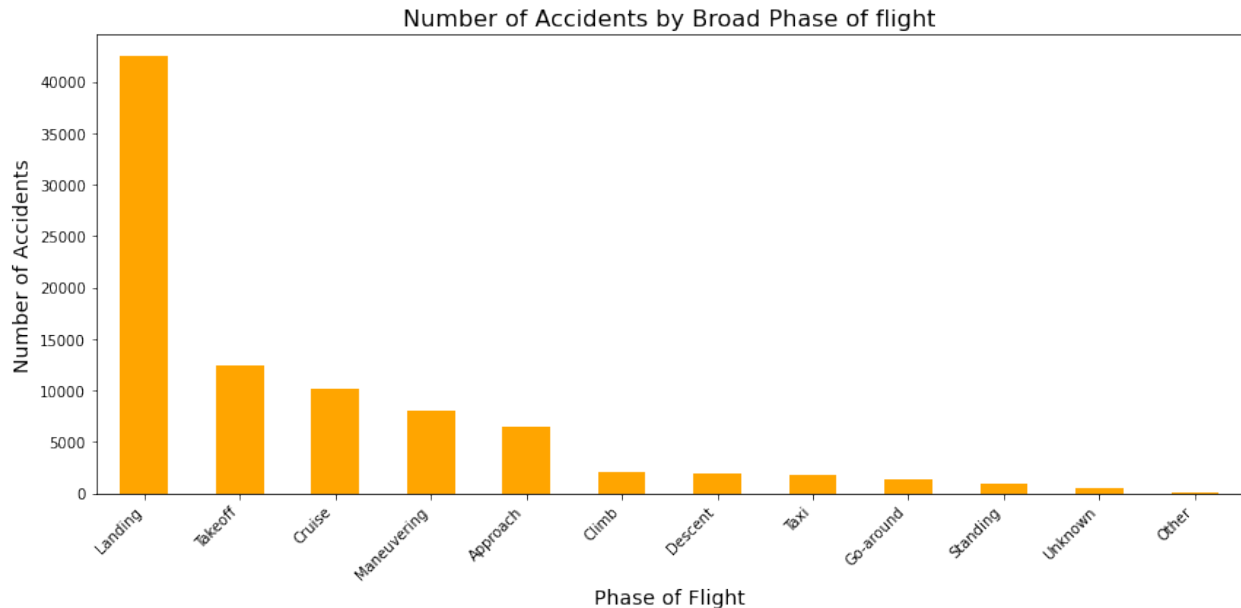
```
# Groups by Broad Phase of Flight and count number of accidents
phase_counts = cleaned_df['Broad.phase.of.flight'].value_counts()

# Plot
fig, ax = plt.subplots(figsize=(12, 6))
phase_counts.plot(kind='bar', ax=ax, color='orange')

#Set titles and labels
ax.set_title('Number of Accidents by Broad Phase of flight',
            fontsize=16)
ax.set_xlabel('Phase of Flight', fontsize=14)
ax.set_ylabel('Number of Accidents', fontsize=14)
ax.set_xticklabels(ax.get_xticklabels(), rotation=45, ha='right')

plt.tight_layout()

#saves image to my device
plt.savefig('Fatality Rates by boeing Aircraft Mode', dpi=300)
plt.show()
```



Most accidents occur during landing !!!

Recommendations.

1. Basing on this analysis, the Two most used models basing on the number of reported incidences are CESSNA and BOEING.
2. CESSNA'S most used model tends to be the 172 but it is not the safest model.Safe ones are the 180,150M,150
3. BOEING'S most used model tends to be 737-200 but it is not the safest model.Safe ones are the 747,A75N1 and 777
4. CESSNA is the most used airplane since 2000 to 2022
5. All the other planes are of low risk because they have less reported cases.

It would therefore be ideal if the company invests in this two airplane because they are the most used brands in the market thereby making them reliable. For short flights that only carry a smaller number of people, CESSNA would do an incredible job. BOEING would be ideal for long commercial flights that carry a good number of people. To avoid accidents, the company should always do regular check ups on both planes specifically on its landing ability and take off because this is when most accidents occur. Employing well experienced pilots would also avoid such incidences. All the other airplanes are also reliable if the company wouldnt want to consider market demand.

Convert my file to csv

```
#I am coverting my file to csv format.  
df.to_csv("Cleaned_df.csv", index=False)
```